Machine Learning
Module 2: Data Exploration

Arnold Vialfont

Master in Management, Business Analytics, HEC UNIL

Spring 2026

Table of Contents

Table of Contents

Data set information

Meta data refers to information (data) about the data set itself. In data science, this usually includes

- Its origin and/or source, if possible, the first one (not only the web page from which it was retieved).
- Its dates of creation and/or retrieval.
- Its name and/or title.
- Its file type (csv, xlxs, etc.).
- Its shape and a general description of its content (each column for tabular data).

Table of Contents

EDA

EDA stands for Exploratory Data Analysis. It is **not** optional! Aims:

- Know the data: e.g., number of modalities for categorical variables, data length, number of missing values, etc.
- Be informed of their general behavior: modality proportions, location and dispersion of numerical variables, etc.
- Detect any outliers or special modes.
- Inspect some of the data relationships.

There exist endless possibilities of EDA. A range of classical and efficient methods is presented below.

EDA strategy

Without prior knowledge of the case:

1. One by one: univariate analysis.
2. By pairs: bivariate analysis.
3. More variables if needed and/or possible.

Exploratory tools

For univariate analysis (one variable):

Graphical:

- Categorical variables: barplot,
- Numerical variables: boxplot, violin plot, and histogram.

Numerical:

- Categorical variables: frequencies and proportions,
- Numerical variables: locations (mean, median, min, max), dispersions (standard deviation, IQR, range), quantiles $(0.25, 0.75)$,
- Both: number of observations, number of missing values.

Exploratory tools

For bivariate analysis (two variables):

Graphical:

- cat*cat: barplots, mosaic plots
- cat*num: boxplots or histograms of num per modality of cat
- num*num: scatterplot

Numerical:

- cat*cat: table of frequencies or proportions
- cat*num: statistics of num per modality of cat
- num*num: correlation

Exploratory tools

For more than two variables this mainly depends on the objective. The complexity increases fast and no universal method exists. Consider

- cat*cat*cat: Sankey diagram.
- num*num*num: parallel coordinates.
- num*num*cat: scatterplot (num*num) with colors or shapes (cat).
- cat*cat*num: can turn num into cat (intervals) or can make cat*cat = cat (combine modalities).

Table of Contents

Numerical variables

Many possibilities. Consider,

- Univariate: x^2, $\ln(x)$, $\ln(1 + x)$, $|x|$, $1/x$, ranks, etc.
- Multivariate: $x_1 x_2$, $x_1 + x_2$, etc.
- Num to ordinal: categorize x by intervals.

Applying any transformation must be guided by the aim of the study (e.g., use $\ln(1 + x)$ on the outcome $x \geq 0$).

Categorical variables

Many possibilities. Consider,

- Turn ordinal into numerical: e.g., "Small", "Medium", "Large" $\rightarrow$ 1, 2, 3.
- Diminish the number of modalities if too many: e.g., There are 10 modalities A, B, . . . , J (100% of observations) but A, B, C covers 80% of the observations $\rightarrow$ create a category "other" and use A, B, C, "other".
- Create dummy variables.

Dummy variables

ML algorithms perform mathematical operations on the data. They require numbers and cannot be performed on categorical data (characters, strings, etc.). These are then represented as **dummy variables**.

Most functions in R handle these representations automatically. However, it is not always the case in Python or in other computer programs. In any cases, it is important to know what is being the scene and how to represent categorical variables as dummy variables.

Dummy variable

Transform a categorical variable into several numerical variables:

- One modality is the **reference** (choice of the user).
- For the other modalities: 0/1 variables, one per modality.
- Value is 1 if the modality is the same as the modality of the dummy variable, 0 otherwise.
- One variable, called **intercept**[1], made of 1's, is added.

Transforming one categorical variable into dummy variables is called **one-hot encoding**.

---

[1]In most algorithm, especially in regressions. In Neural Network, it is called the **bias**.

Dummy variable example

A variable indicates a shape. It has three levels Cloud, Moon, and Star. The reference level below is Cloud. Below, 15 instances.

Levels
Cloud
Star
Moon

Shape
Cloud
Cloud
Star
Cloud
Moon
Cloud
Moon
Moon
Star
Star
Star
Star
Moon
Cloud
Moon

Reference: Cloud		
Intercept	Shape_Star	Shape_Moon
1	0	0
1	0	0
1	1	0
1	0	0
1	0	1
1	0	0
1	0	1
1	0	1
1	1	0
1	1	0
1	1	0
1	1	0
1	0	1
1	0	0
1	0	1

Dummy variable example

The same example when the reference level below is Moon.

Levels		Shape		Intercept	Shape_Cloud	Shape_Star
					Reference: Moon	
Cloud		Cloud		1	1	0
Star		Cloud		1	1	0
Moon		Star		1	0	1
		Cloud		1	1	0
		Moon		1	0	0
		Cloud		1	1	0
		Moon		1	0	0
		Moon		1	0	0
		Star		1	0	1
		Star		1	0	1
		Star		1	0	1
		Star		1	0	1
		Moon		1	0	0
		Cloud		1	1	0
		Moon		1	0	0

Table of Contents

Dealing with (without) missing values

Missing value analysis is **never** easy and would deserve a whole course. In a quick fix, consider

- Detect: which variables, which code (e.g., NA, -999, etc.).
- Quantify: how many (number and proportion) for variables and cases.
- Relate: two (or more) variables are systematically missing together?

In absence of data, lots of techniques will break down. Can this be solved perfectly? No.

Dealing with (without) missing values

Various techniques, none of which is perfect:

- Remove the cases with at least one missing feature.
- Remove the feature with too many missing values.
- Input the missing value.

Remove cases

Pro:

- Very simple and systematic

Cons:

- Can leave no data: e.g., one variable is 99% systematically missing.
- Can bias the analysis: e.g., missing is linked to the outcome (NMAR)

Remove features

Pro:

- if most of the missing values are concentrated in one feature, then it saves lot of cases.

Cons:

- Can leave out important information.

Consider also replacing the feature by indicator "0/1" for missing or not.

Imputation

- **Naive:** per feature, replace NA by average or median (num), or the most frequent modality (cat).
  - Pro: None.
  - Cons: Avoid if possible (even if it is often used).
- **Model based:** use the other features to predict the missing one (case by case).
  - Pro: Use all the available information.
  - Cons: Incorrectly diminish the noise in the data which enforce the correlation and give a wrong impression of certainty.

More advanced methods: repeated random imputation, EM algorithm, etc.

Table of Contents

Definition

Behind the term **outlier** hides lots of ideas. It can be (non-exhaustive)

- An impossible value: exceeding physical limits.
- An extreme value of a feature: possible but out of some predefined bounds.
- An unlikely value: unlikely given the other features.

The second case is limited to numerical features. The two others can be categorical features.

Detection

For the three previous cases, consider

- Often easy to detect from prior information. E.g., negative surface in a real estate data study, time exceeding 1 hour in a 1-hour observation study, etc.
- Use boxplots to point them: values beyond 1.5 IQR from the quartile. E.g., age of 101 years when the other largest observed is 35 years.
- Complex: use prior knowledge or a model. E.g., a room used for an RMI in a hospital for a specific diagnostic for which this room is never used, a yearly wage greater than \$1′000′000 for a job in manufacturing.

Discussion

With outliers, things are never simple and automatic removal is often a bad idea. In an ideal world, only errors should be removed. Consider,

- **Check the facts**: how many (if only one then it could be removed safely), where, which features?
- Inquiry **how the data were gathered**: e.g., 5% of the revenues are negative. Did you consider possible that the revenue could indeed be negative in that study?
- If set aside then **document it**. This should never be hidden.
- In some situation, outliers are of interest: e.g., fraud analytics.

Table of Contents

Further words

- Univariate: be exhaustive, analyze all variables systematically.
- Bivariate: select the interesting ones (e.g., in supervised learning, outcome vs other).
- More than bivariate: be selective.

General advices:

- Avoid using too many modalities in a categorical variable.
- Do not compute correlations between dummy variables.
- Do not use missing value indicator as a number (e.g., -999).
- Transformations are not neutral. They influence the final conclusion. They must be taken into account and transparent.
- Think that EDA is not a one way through. It is a recursive/trial-and-error process (not unusual 95% of the work...)