

# Machine Learning

## Module 3.1 - Models: Linear and Logistic regressions

Arnold Vialfont

Master in Management, Business Analytics, HEC UNIL

Spring 2026

# Table of Contents

- 1 Linear regression
- 2 Logistic regression
- 3 Interpretation
- 4 Selection of variables

# Table of Contents

1 Linear regression

2 Logistic regression

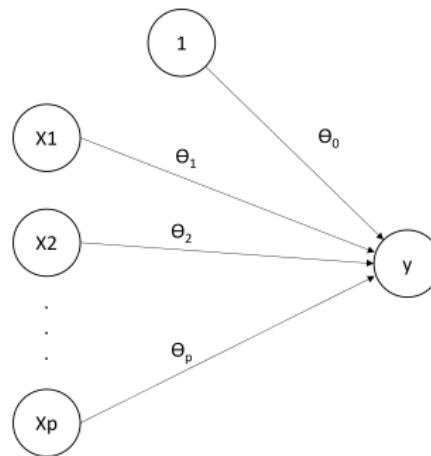
3 Interpretation

4 Selection of variables

# Linear regression

The prediction formula is

$$f(x; \theta) = \theta_0 + \theta_1 x_1 + \dots + \theta_p x_p$$



# The MSE and the OLS

The loss function is the **mean squared error**, MSE,

$$\bar{\mathcal{L}}(\theta) = \frac{1}{n} \sum_{i=1}^n \{y_i - f(x_i; \theta)\}^2.$$

The optimal parameter under this loss is called the **Ordinary Least Squares**, OLS,

$$\hat{\theta} = \arg \min_{\theta} \frac{1}{n} \sum_{i=1}^n \{y_i - f(x_i; \theta)\}^2.$$

# The OLS

To find the OLS, the algorithm is exact and the final solution can be explicitly computed with the matrix-operation

$$\hat{\theta} = (X^T X)^{-1} X^T y,$$

where

- $X$  is the so-called **design matrix** of size  $n \times (p + 1)$  and whose  $i$ -th row contains

$$[1, x_{i1}, \dots, x_{ip}],$$

- $y$  is the vector of length  $n$  whose  $i$ -th element is  $y_i$ .

# Table of Contents

1 Linear regression

2 Logistic regression

3 Interpretation

4 Selection of variables

# Logistic regression

The logistic regression is a model for **binary classification**. Prediction formula is in several steps:

- Compute the **linear predictor**

$$z(x; \theta) = \theta_0 + \theta_1 x_1 + \cdots + \theta_p x_p,$$

- Compute the **probability prediction** (sigmoid function):

$$p(x; \theta) = P(Y = 1 | X = x) = \frac{\exp\{z(x; \theta)\}}{1 + \exp\{z(x; \theta)\}}.$$

- Compute the **prediction of the class**:

$$f(x; \theta) = \begin{cases} 1, & \text{if } p \geq 0.5, \\ 0, & \text{if } p < 0.5. \end{cases}$$

# The sigmoid and the logit function

The **logit function** is the inverse of the **sigmoid function** and thus transforms  $p(x; \theta)$  to  $z(x; \theta)$ .

$$z(x; \theta) = \log \frac{p(x; \theta)}{1 - p(x; \theta)}$$

# Estimation

The loss function uses the probabilities  $p(x; \theta)$  and not the final predictions. The loss is the **cross-entropy** also called **negative log-likelihood**:

$$\mathcal{L}(y, p) = -y \log p - (1 - y) \log(1 - p).$$

Interpretation,

- If  $y = 1$ , we want  $p$  to be large (close to 1). The contribution to the loss is

$$\mathcal{L}(1, p) = -\log p$$

It will be small indeed if  $p$  is large.

- If  $y = 0$ , we want  $p$  to be small (close to 0). The contribution to the loss is

$$\mathcal{L}(0, p) = -\log(1 - p)$$

It will be small indeed if  $p$  is small.

# Estimation

The overall loss is

$$\bar{\mathcal{L}}(\theta) = -\frac{1}{n} \sum_{i=1}^n y_i \log p(x_i; \theta) + (1 - y_i) \log \{1 - p(x_i; \theta)\}.$$

The log in this formula can be in any base. Often,

- Machine learners use  $\log_2$  for binary classification,
- Statisticians use  $\ln$ .

This has no consequence on the final result. But it can bring confusion from time to time.

# Optimal parameters

To obtain the optimal parameters, the best algorithm is the **Newton-Raphson** algorithm. It requires

- To compute the first and second derivatives of  $\bar{\mathcal{L}}(\theta)$ ,
- To build a sequence of  $\hat{\theta}_k$  that converges to the optimal one using these derivatives.

This algorithm is very fast and efficient. However, there is no explicit formula for  $\hat{\theta}$ , unlike the OLS.

The optimal  $\hat{\theta}$  is called the **maximum likelihood estimator, MLE**.

# Table of Contents

1 Linear regression

2 Logistic regression

3 Interpretation

4 Selection of variables

# Interpretation

Linear and logistic regressions are highly interpretable in that

- the coefficients quantify the link between the features  $x$  and the outcome  $y$ .
- the certainty of the prediction can be quantify.

# Interpretation of the coefficients

For the **linear regression**, coefficients are interpreted as *slopes*.

- When the feature  $x_1$  increases by 1 unit, the outcome  $y$  increases in average by  $\theta_1$  units (same for all features  $1, \dots, p$ ).
- A positive coefficient  $\theta_j$  means a positive linear association between the feature  $x_j$  and the outcome  $y$ .
- The larger the coefficients, the larger the association, in absolute value. (note: pay attention to the scale!)
- For the categorical features, the coefficients estimate the average change in the outcome  $y$  when the feature switches from the reference level to any other level. It is thus a **contrast** with the reference level.

Note: one should not say that an increase in  $x_j$  *causes* an increase of the response. It is an association. The causality implies a direction and is more complex to establish.

# Interpretation of the coefficients

For the **logistic regression**, because of the sigmoid transform, the interpretation of the coefficients is not direct like with the linear regression:

- With a positive  $\theta_j$ , an increase of  $x_j$  is associated with an increase of the probability that  $y = 1$ .
- The larger the coefficient, the larger the increase. However, it the increase is not linear and depends on the other features.
- A negative coefficient means a decrease in the probability of the positive class ( $y = 1$ ).

Note: coefficients can be interpreted in terms of odds ratio. We will not pursue this here.

# Certainty for linear regression

For **linear regression**, certainty can be measured by **prediction intervals**. In practice, the main interest relies in the prediction interval for **the future value**<sup>1</sup>.

Let  $x$  be a set of new features, the **point** prediction for  $y(x)$  uses the estimate  $\hat{\theta}$ :

$$f(x; \hat{\theta}) = \hat{\theta}_0 + \hat{\theta}_1 x_1 + \cdots + \hat{\theta}_p x_p.$$

Now, rather than a point estimate, we want to build an interval  $[L, U]$  such that

$$P(L \leq y(x) \leq U) = 1 - \alpha,$$

where  $\alpha$  is usually set to 5% for an interval at 95%. To build this interval, we rely on probabilistic assumptions of the model.

---

<sup>1</sup>As opposed to prediction interval for the mean.

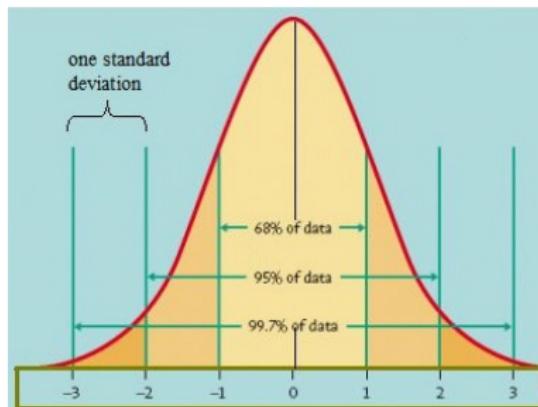
# Certainty for linear regression

It is often assumed that the true response  $y(x)$  for feature  $x$  satisfies

$$y(x) = \theta_0 + \theta_1 x_1 + \cdots + \theta_p x_p + \sigma e = f(x; \theta) + \sigma e,$$

where the residual  $e$  is standard normal,  $e \sim N(0, 1)$ , and  $\sigma$  is the standard deviation.

In particular, we expect the residual distribution to be symmetric around 0 and (informally) respect the 68-95-99.7 rule:



# Certainty for linear regression

Therefore, to find  $[L, U]$ , we set

$$0.95 = P(L \leq y(x) \leq U) = P\left(\frac{L - f(x; \theta)}{\sigma} \leq e \leq \frac{U - f(x; \theta)}{\sigma}\right).$$

Using the normal distribution, we find that

$$L - f(x; \theta) = -\sigma z_{1-\alpha/2}, \quad U - f(x; \theta) = \sigma z_{1-\alpha/2}.$$

This gives

$$L = f(x; \theta) - \sigma z_{1-\alpha/2}, \quad U = f(x; \theta) + \sigma z_{1-\alpha/2}.$$

For  $\alpha = 5\%$ ,  $z_{1-\alpha/2} = 1.96$ , and

$$[L, U] = [f(x; \theta) \pm 1.96\sigma].$$

# Certainty for linear regression

Using a plug-in estimate, this gives us a *rough* prediction interval (at 95%) of

$$[\hat{L}, \hat{U}] = [f(x; \hat{\theta}) \pm 1.96s],$$

where  $s$  is the unbiased estimate of the standard deviation

$$s^2 = \frac{1}{n - (p + 1)} \sum_{i=1}^n \{y_i - f(x_i; \hat{\theta})\}^2.$$

This prediction interval is however hardly used/implemented because

- the estimate  $s$  carries uncertainty. If taken into account, the Student  $t_{n-(p+1)}$  distribution should be used.
- the estimate  $\hat{\theta}$  carries uncertainty. If taken into account, the estimate of  $s$  should be changed to  $s\sqrt{1 + x^T(X^T X)^{-1}x}$ .

Both adaptations widen the interval.

# Certainty for logistic regression

The probability provides an interpretation of the certainty the model provides on a classification. Consider:

- $\hat{y} = 1$  with  $\hat{p} = 0.99$ : the prediction is **certain**.
- $\hat{y} = 1$  with  $\hat{p} = 0.53$ : the prediction is **uncertain**.

In both cases, the predicted class is the same but the probability provides a more precise view on it: if the instance is far in the class or on the edge between the two classes.

The prediction "Good" for a customer with a probability of 0.51 is uncertain. Alternatively, the prediction rule could set to a larger value than 0.5 to increase the certainty.

Also, a model with a lot of predictions close to 0.5 is considered as poor because of its uncertainty<sup>2</sup>.

---

<sup>2</sup>A little thought on that will bring us back to the entropy.

# Table of Contents

- 1 Linear regression
- 2 Logistic regression
- 3 Interpretation
- 4 Selection of variables

# Occam's Razor

A parsimony principle:

"All else being equal, the simplest explanation is the best one"

- In other words, among two models with approximately the same prediction quality, choose the simplest one.
- In practice, we remove from the model variables that do not impair too much the prediction quality.

Simplifying the model is a solution to **overfitting**. This will be studied later in the course.

# Akaike Information Criterion

The **AIC** (Akaike Information Criterion) is

$$AIC = -2\hat{\ell} + 2k$$

where

- $\hat{\ell}$  is the maximum **log-likelihood** and measure the goodness-of-fit,
- $k$  is the number of parameters and measure the model complexity.

**Minimizing** the AIC achieves a trade-off between the quality of prediction and the model complexity.

# Akaike Information Criterion

For linear regressions,

- The number of parameters is  $k = p + 2$  with  $\theta_0, \theta_1, \dots, \theta_p$  and  $\sigma$ ,
- The log-likelihood part equals

$$-2\hat{\ell} = n \ln 2\pi + n \ln \hat{\sigma}^2 + \frac{1}{\hat{\sigma}^2} \sum_{i=1}^n \left\{ y_i - f(x; \hat{\theta}) \right\}^2,$$

where  $\hat{\sigma}^2 = (n - p - 1)s^2/n$ .

For logistic regressions,

- The number of parameters is  $k = p + 1$  for  $\theta_0, \theta_1, \dots, \theta_p$ ,
- The log-likelihood part equals

$$-2\hat{\ell} = 2 \sum_{i=1}^n y_i \ln p(x_i; \hat{\theta}) + (1 - y_i) \ln \{1 - p(x_i; \hat{\theta})\}.$$

# Variables selection with AIC

Automatic variable selection using stepwise minimization of the AIC can be performed. There are

- Backward: start from the most complete model and try to remove variable one at a time (if it decreases the AIC)
- Forward: start from the empty model and try to add one variable at a time (if it decreases the AIC).
- Both: start to add or remove at each step.

At each step, all the models in competition are fitted. The procedure is computationally intense.

# Variable selection with penalization

A different approach consists of penalizing the loss function so that, during the training of the parameters, the variable selection applies directly.

The most common penalties are:

- $L_1$  penalty (LASSO)

$$\min_{\theta} \bar{\mathcal{L}}(\theta) + \lambda \sum_{j=1}^p |\theta_j|$$

- $L_2$  penalty (Ridge)

$$\min_{\theta} \bar{\mathcal{L}}(\theta) + \lambda \sum_{j=1}^p \theta_j^2$$

Usually,  $\theta_0$  is not penalized.

# Variable selection with penalization

The penalty parameter  $\lambda \geq 0$ :

- If  $\lambda = 0$ , then there is no penalty.
- If  $\lambda \rightarrow \infty$ , then  $\theta \rightarrow 0$ .

For intermediate values, some components of  $\theta$  will be small, pushed toward 0.

This is equivalent to variable selection: setting  $\theta_j = 0$  is equivalent to not including  $x_j$ .

Selection of  $\lambda$  can be done with **cross-validation** (CV ; see later).

# Variable selection with penalization

- $L_1$  shrink some of the  $\theta_j$ , set some  $\theta_j = 0$ , select variables.
- $L_2$  shrink all the  $\theta_j$ 's, avoiding extremes  $\theta_j$ , regularize  $\theta$ .

**Elastic net** combines  $L_1$  and  $L_2$ :

$$\min_{\theta} \bar{\mathcal{L}}(\theta) + \lambda \left\{ \alpha \sum_{j=1}^p |\theta_j| + \frac{1-\alpha}{2} \sum_{j=1}^p \theta_j^2 \right\}$$

with  $0 \leq \alpha \leq 1$ ,

- If  $\alpha = 0$ , it is the ridge ( $L_2$ )
- If  $\alpha = 1$ , it is the LASSO ( $L_1$ )

Often,  $\lambda$  is selected by the data (CV), while  $\alpha$  is set by the user.